

**Author: Suresh Lokiah**

**LinkedIn: [linkedin.com/in/sureshlokiah](https://www.linkedin.com/in/sureshlokiah)**

**Email: ([sureshlokiah@gmail.com](mailto:sureshlokiah@gmail.com))**



( Image Credits: <https://www.courier-journal.com/story/news/politics/metro-government/2019/06/20/louisville-bird-lime-may-be-joined-5-dockless-scooters-lyft-jump-spin-bolt/1499757001/> )

## Introduction

This exercise introduces the concepts of Data Wrangling and Machine Learning using dockless(scooter and bike) *Trips* data from Louisville, KY

<https://data.louisvilleky.gov/dataset/dockless-vehicles>. The datafile used here is downloaded from the link showing 'Dockless Trips 02/01/2020-07/08/2020'.

The data.louisvilleky.gov website also provides the data definition for the fields listed in the data.

### Stages

The exercise is organized (roughly) in these stages

1. Data Handling
2. Perform EDA
3. Cleanup
4. Visualizing the Data (Numeric and Spatial data)
5. Machine Learning - doing Model selection

```
In [1]: ## Install and Update python packages

#!pip install pandas
#!pip install seaborn
#!pip install sklearn

#!pip3 list --outdated --format=freeze | grep -v '^\-e' | cut -d = -f 1 | xargs

In [2]: ## Import

import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

import numpy as np
import seaborn as sns

## Check if we have the latest version of Pandas, if not lets upgrade
pd.__version__

Out[2]: '1.4.3'
```

## Stages

### Data Handling

```
In [3]: ### Open Google Drive and download 'Louisville-Dockless-Trips.csv'

#from google.colab import drive
#drive.mount("/content/gdrive")
#df = pd.read_csv('/content/gdrive/My Drive/Louisville-Dockless-Trips.csv')

In [4]: # Load Dataset
df = pd.read_csv('/Users/suresh/Downloads/Louisville-Dockless-Trips.csv')
```

### Perform EDA

```
In [5]: ## Find the count of (Rows, Columns)

df.shape
```

Out[5]: (32962, 13)

So the dataset is made up of **32962 Rows and 13 Columns**

```
In [6]: # Checkout the Rows and Columns and DataTypes

df.info();
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32962 entries, 0 to 32961
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   TripID                 32962 non-null  object
1   StartDate              32962 non-null  object
2   StartTime              32962 non-null  object
3   EndDate                32962 non-null  object
4   EndTime                32962 non-null  object
5   TripDuration           32962 non-null  int64
6   TripDistance           32962 non-null  float64
7   StartLatitude          32962 non-null  float64
8   StartLongitude         32962 non-null  float64
9   EndLatitude            32962 non-null  float64
10  EndLongitude           32962 non-null  float64
11  DayOfWeek              32962 non-null  int64
12  HourNum                32962 non-null  int64
dtypes: float64(5), int64(3), object(5)
memory usage: 3.3+ MB
```

In [7]: *## Review the first few rows for data in column*

```
df.head()
```

Out[7]:

	TripID	StartDate	StartTime	EndDate	EndTime	TripDuration	TripDistance	StartLat
--	--------	-----------	-----------	---------	---------	--------------	--------------	----------

0	dd62d00f-11cf-5a7b-b6df-a1917e90eb09	2/1/2020	8:15	2/1/2020	8:30	14	1.358	3
1	ea963a4b-b2b9-5ab9-8442-cbff3878e32c	2/1/2020	8:30	2/1/2020	8:45	6	0.195	3
2	d081f7b8-5d6f-54a8-adf1-fa2c931b8440	2/1/2020	8:45	2/1/2020	8:45	3	0.118	3
3	135100da-29e7-5e7b-b150-73784e2d983f	2/1/2020	9:00	2/1/2020	9:00	3	0.252	3
4	2fea57ae-fb0e-5538-85a3-749d4854cfc	2/1/2020	9:15	2/1/2020	9:30	16	0.434	3

In [8]: *# Convert to datetime[ns] field*

```
df.StartDate = df.StartDate.astype('datetime64[ns]')
df.EndDate = df.EndDate.astype('datetime64[ns]')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32962 entries, 0 to 32961
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   TripID                 32962 non-null  object
1   StartDate              32962 non-null  datetime64[ns]
2   StartTime              32962 non-null  object
3   EndDate                32962 non-null  datetime64[ns]
4   EndTime                32962 non-null  object
5   TripDuration           32962 non-null  int64
6   TripDistance           32962 non-null  float64
7   StartLatitude          32962 non-null  float64
8   StartLongitude         32962 non-null  float64
9   EndLatitude            32962 non-null  float64
10  EndLongitude           32962 non-null  float64
11  DayOfWeek              32962 non-null  int64
12  HourNum                32962 non-null  int64
dtypes: datetime64[ns](2), float64(5), int64(3), object(3)
memory usage: 3.3+ MB
```

```
In [9]: ## Check the date range

print("Range of Start Date = " + str(df.StartDate.min()) + " - " + str(df.StartDate.max()))

Range of Start Date = 2020-02-01 00:00:00 - 2020-07-08 00:00:00
```

```
In [10]: #Describe the data

df.describe()
```

```
Out[10]:
```

	TripDuration	TripDistance	StartLatitude	StartLongitude	EndLatitude	EndLongitude
count	32962.000000	32962.000000	32962.000000	32962.000000	32962.000000	32962.000000
mean	21.033099	1.096729	38.244277	-85.738711	38.244402	-85.738500
std	24.094454	4.213029	0.098125	0.777303	0.098135	0.777200
min	0.000000	0.000000	33.437000	-121.987000	33.423000	-121.895000
25%	5.000000	0.092000	38.230000	-85.757000	38.230000	-85.757000
50%	13.000000	0.543000	38.251000	-85.745000	38.251000	-85.745000
75%	29.000000	1.339000	38.258000	-85.735000	38.258000	-85.736000
max	1013.000000	596.139000	50.110000	8.674000	50.117000	8.685000

### TripDuration

- mean: 21.033099 and std: 24.094454 - Shows (likely) issues with data because the difference of (mean - std) would be -ve value
- mean: 21.033099 and 50%: 13.000000 - indicates that the data is right skewed.
- max: 1013.000000 - is likely a outlier or bad data
- min: 0.000000 - has 0's for Trip Duration that must be dropped.

### TripDistance

- min: 0.000000 - has 0s for Trip Distance that must be dropped

- max: 596.139000 - is likely a outlier or bad data

## DayOfWeek

- min: 1 and max: 7 representing 7 days, and 1 being 'Sunday' from the data definition - <https://data.louisvilleky.gov/dataset/dockless-vehicles>.

## HourNum

- min: 6.000000 and max: 23.000000 indicating data between 06:00 to 23:00 (and not 23:01 or 23:05...)

# Cleanup

1. Drop duplicate rows (if any)
2. Drop rows having 0's in Trip Duration or Trip Distance
3. Create new Field DayOfWeekStr from DayOfWeek field

## Drop duplicate rows

```
In [11]: # Check if there is any duplicate rows;
# Zero would indicate no duplicates.

print(len(df) - len(df.drop_duplicates()))

0
```

## Drop rows having 0's

```
In [12]: # Count the rows having TripDistance == 0 and TripDuration == 0

print(len(df[(df.TripDistance == 0) & (df.TripDuration == 0)]))

72
```

```
In [13]: ## Drop rows having 0's in the TripDistance and TripDuration

df = df[(df.TripDistance != 0) & (df.TripDuration != 0)]
```

```
In [14]: df.shape
```

```
Out[14]: (31847, 13)
```

## Create new Field

```
In [15]: ## Convert DayOfWeek Number to String and save it in a new field DayOfWeekStr

d_to_week = {
    1: 'Sunday',
    2: 'Monday',
    3: 'Tuesday',
```

```

4: 'Wednesday',
5: 'Thursday',
6: 'Friday',
7: 'Saturday'
}

```

```

df['DayOfWeekStr'] = df.DayOfWeek.replace(to_replace=d_to_week)
df.DayOfWeekStr.head()

```

```

Out[15]: 0    Saturday
1    Saturday
2    Saturday
3    Saturday
4    Saturday
Name: DayOfWeekStr, dtype: object

```

## Visualizing the Data

1. Graph discrete data( HourNum, DayOfWeek)
2. Graph continuous data(TripDuration, TripDistance)
3. Graph the Geo data (Lat, Long)

### Graph discrete data


#### Plot Counts by HourNum

```

In [16]: ## Plot Counts by HourNum using sns countplot()
         #df.HourNum.value_counts()

sns.countplot(data=df, x='HourNum',color='b');
plt.title('Counts by HourNum');

```

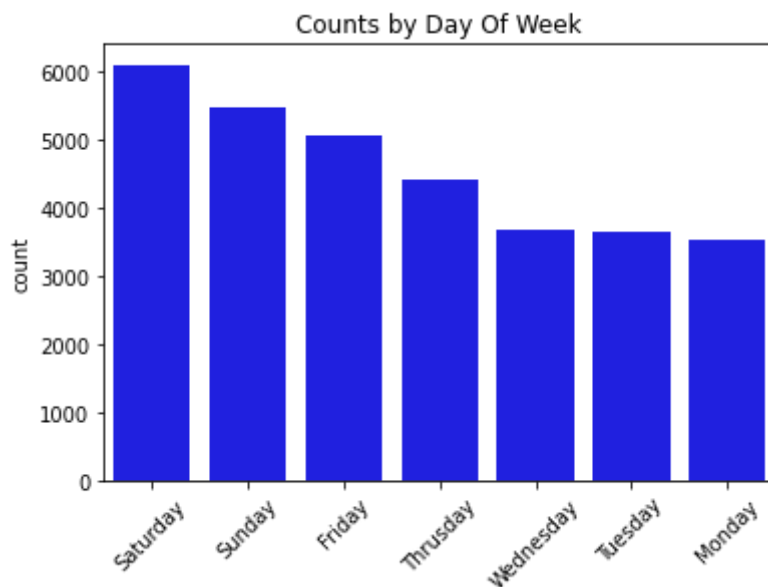


Shows Peak usage at 8:00 pm during the day. The other higher usages are at 9:00pm and 10:00 pm. And, as expected the usage ramps-up starting at 6:00 am.

#### Plot Counts by DayOfWeek

```
In [17]: ## Bar Plot Counts by Day Of Week

sns.barplot(data=df, x=df.DayOfWeekStr.value_counts().index, y=df.DayOfWeek.val
plt.title('Counts by Day Of Week');
plt.xticks(rotation = 45);
plt.ylabel('count');
```



Weekends have the highest usage (Saturday, Sunday, Friday) and Monday has the relatively lower usage.

## Subplot continuous variables

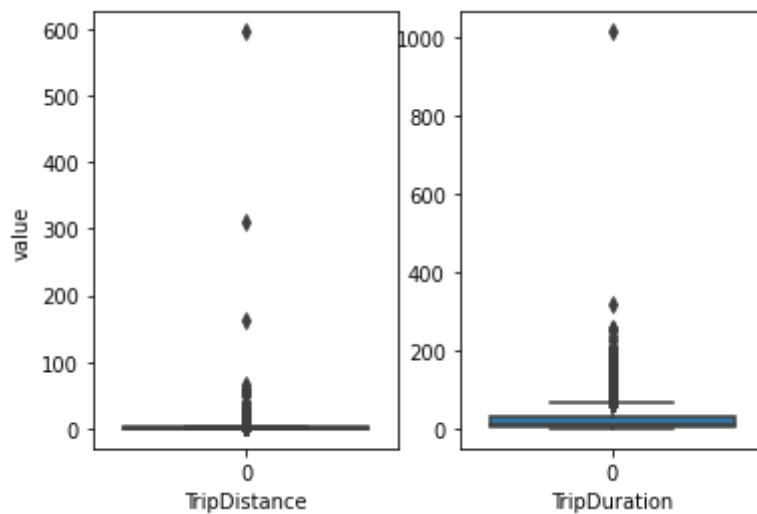
```
In [18]: #Plot the continuous values TripDuration, TripDistance

fig, ((ax11, ax12)) = plt.subplots(ncols=2)

#plt.title("Box Plots for Trip Distance and Duration")

sns.boxplot(data=df.TripDistance, ax=ax11);
ax11.set_xlabel('TripDistance')
ax11.set_ylabel('value');

sns.boxplot(data=df.TripDuration, ax=ax12);
ax12.set_xlabel('TripDuration')
ax11.set_ylabel('value');
```



There is lots of outliers. Lets remove the outliers by  $1.5 \times \text{IQR}$  for TripDistance and TripDuration

## Outlier removal

```
In [19]: ### Method to subset data by iqr to eliminate outliers.

def subset_by_iqr(df, column, whisker_width=1.5):
    """Remove outliers from a dataframe by column, including optional
    whiskers, removing rows for which the column value are
    less than  $Q1 - 1.5 \times \text{IQR}$  or greater than  $Q3 + 1.5 \times \text{IQR}$ .
    Args:
        df (:obj:pd.DataFrame): A pandas dataframe to subset
        column (str): Name of the column to calculate the subset from.
        whisker_width (float): Optional, loosen the IQR filter by a
            factor of `whisker_width` * IQR.

    Returns:
        (:obj:pd.DataFrame): Filtered dataframe
    """
    # Calculate Q1, Q2 and IQR
    q1 = df[column].quantile(0.25)
    q3 = df[column].quantile(0.75)
    iqr = q3 - q1
    # Apply filter with respect to IQR, including optional whiskers
    filter = (df[column] >= q1 - whisker_width*iqr) & (df[column] <= q3 + whisker_width*iqr)
    return df.loc[filter]
```

```
In [20]: # Example for whiskers = 1.5, as requested by the OP

df_filtered = subset_by_iqr(df, 'TripDistance', whisker_width=1.5)
df = df_filtered
print("Count of Rows and Columns after removing the outlier data rows based on TripDistance")

df_filtered = subset_by_iqr(df, 'TripDuration', whisker_width=1.5)
df = df_filtered
print("Count of Rows and Columns after removing the outlier data rows based on TripDuration")
```

```
Count of Rows and Columns after removing the outlier data rows based on TripDistance = (29624, 14)
Count of Rows and Columns after removing the outlier data rows based on TripDuration = (27795, 14)
```




```
In [21]: # Now, Repeat the Box Plots
## Plot the continuous values TripDuration, TripDistance

fig, ((ax11, ax12)) = plt.subplots(ncols=2)

sns.boxplot(data=df.TripDistance, ax=ax11);
ax11.set_xlabel('TripDistance')
ax11.set_ylabel('value');

sns.boxplot(data=df.TripDuration, ax=ax12);
ax12.set_xlabel('TripDuration')
ax11.set_ylabel('value');
```



## Correleation graphs


```
In [22]: # pearson coefficient is a correlation indicator
# generate a pearson coefficient for each peer of variables
pearson = df.corr(method='pearson')
pearson
```

```
Out[22]:
```

	TripDuration	TripDistance	StartLatitude	StartLongitude	EndLatitude	EndLong
TripDuration	1.000000	0.279317	0.054945	0.003374	0.057773	0.00
TripDistance	0.279317	1.000000	0.009922	0.007005	0.011269	0.0
StartLatitude	0.054945	0.009922	1.000000	0.984043	0.998274	0.98
StartLongitude	0.003374	0.007005	0.984043	1.000000	0.983593	0.99
EndLatitude	0.057773	0.011269	0.998274	0.983593	1.000000	0.98
EndLongitude	0.003522	0.007211	0.984030	0.999954	0.983551	1.00
DayOfWeek	0.012513	-0.004071	0.012118	0.004278	0.012297	0.00
HourNum	0.249565	0.015948	0.044947	-0.017455	0.045539	-0.0

```
In [23]: ## Lets do correlation using heatmap

plt.figure(figsize= [10,10])
sns.heatmap(data=df.corr(), annot = True, fmt = '.4f', cmap = 'Paired', center=
```



Insights

Though not significant

Positive Correlation -

- 1. Trip Duration to Trip Distance
- 2. Trip Duration to Hour Num


Negative Correlation -

- 1. HourNum to StartLongitude
- 2. HourNum to EndLongitude
- 3. DayOfWeek to TripDistance

## Sample and Visualize

```
In [24]: ## SubSample data
# random selection of 1000 data points
sample = df.sample(n=1000)
```

```
In [25]: # visualise variable distribution using jointplot
sns.jointplot(data=df, x=sample.TripDistance, y=sample.TripDuration, palette =
```



## Spatial Data

```
In [26]: #!pip install folium
```

```
In [27]: # visualise the spatial data points on maps

import folium
from folium.plugins import HeatMap
from folium.plugins import HeatMapWithTime
```

```
In [28]: # Louisville, KY coordinates

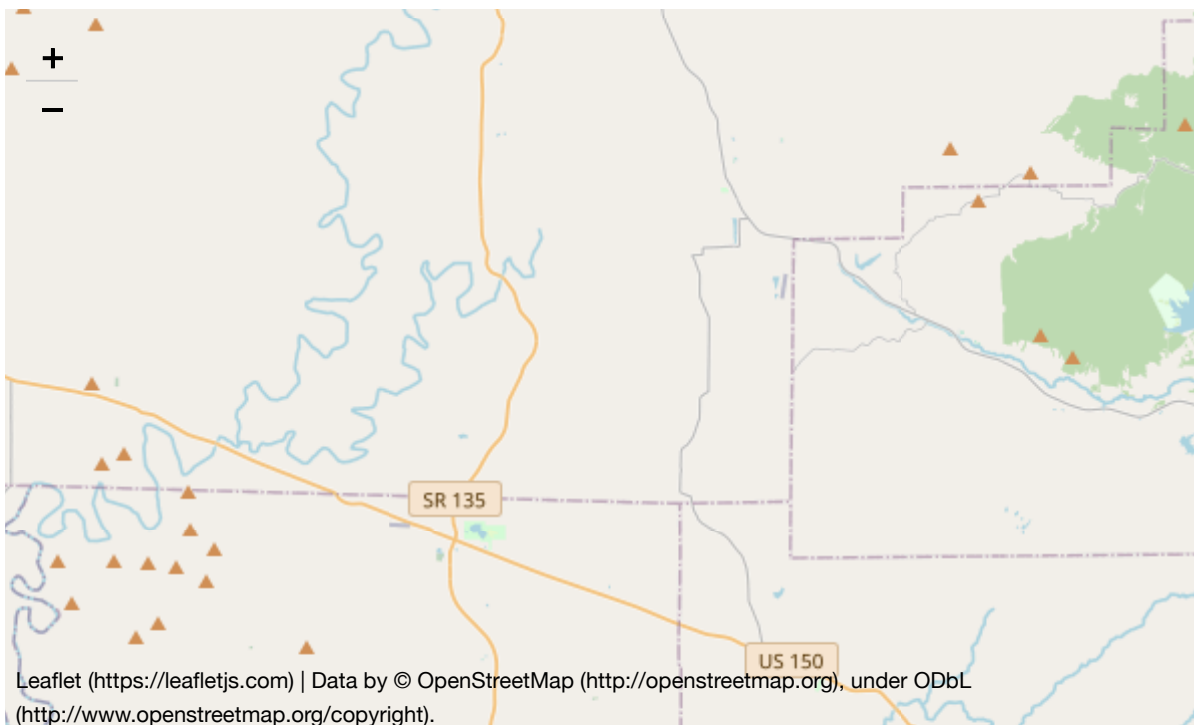
Louisville=[38.2527,-85.7585]
```

## Map of Louisville,KY

```
In [29]: # map of Louisville, KY
map_louisville = folium.Map(location=Louisville,
```

```
map_louisville zoom_start=12)
```

Out[29]:



## Map of Start Latitude, Longitude

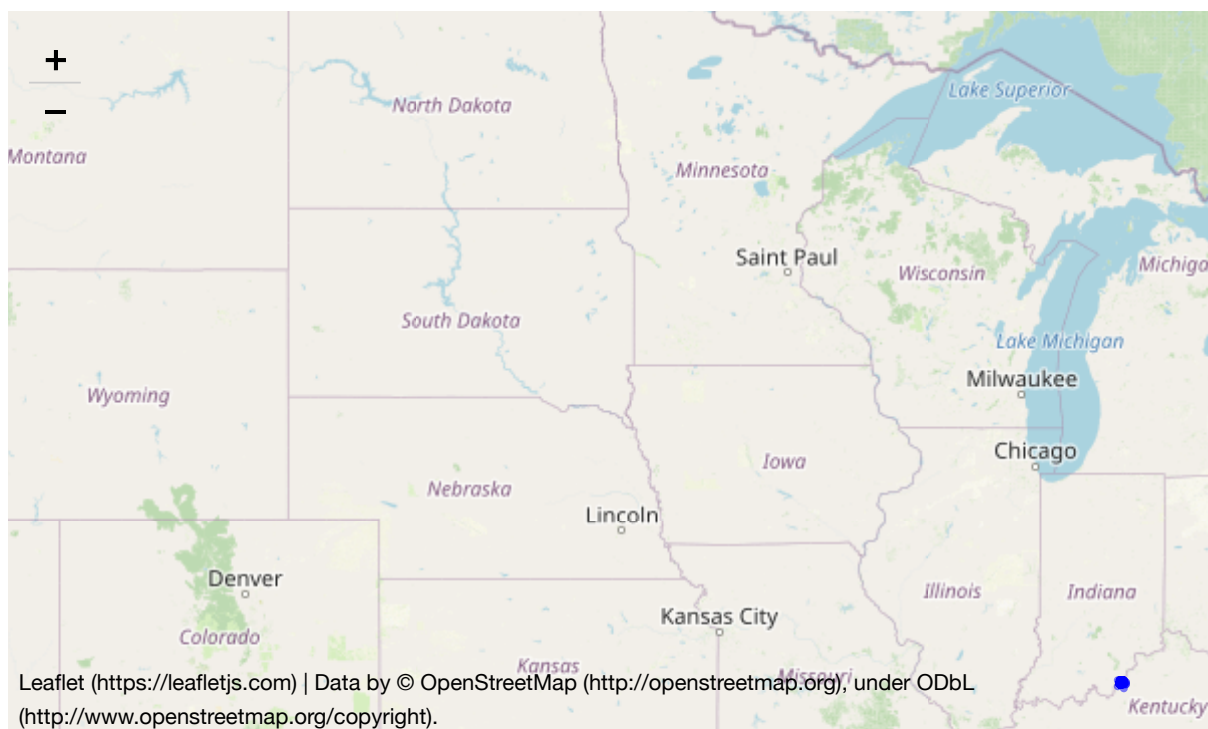
```
In [30]: # map origin points (Start Latitude, and Start Longitude)

map_origin = folium.Map(location=Louisville,
                          zoom_start=12)

for lat, lng in zip(sample['StartLatitude'],
                    sample['StartLongitude']):
    folium.CircleMarker([lat, lng],
                        radius=2,
                        stroke=False,
                        fill_color='blue',
                        fill_opacity=0.5).add_to(map_origin)

map_origin
```

Out [30]:



## Cluster Map of Trips

```
In [31]: # cluster points to make the map more readable
map_cluster = folium.Map(location=Louisville,
                           zoom_start=12)

cluster = folium.plugins.MarkerCluster().add_to(map_cluster)

for lat, lng in zip(sample['StartLatitude'],
                     sample['StartLongitude']):
    folium.Marker([lat, lng], tooltip=[lat, lng]).add_to(cluster)

map_cluster
```

Out [31]:

In [32]: *## Repeat with different Sample, See the concentration of activity*

```

sample2 = df.sample(n=1000)
# cluster points to make the map more readable
map_cluster = folium.Map(location=Louisville,
                          zoom_start=12)

cluster = folium.plugins.MarkerCluster().add_to(map_cluster)

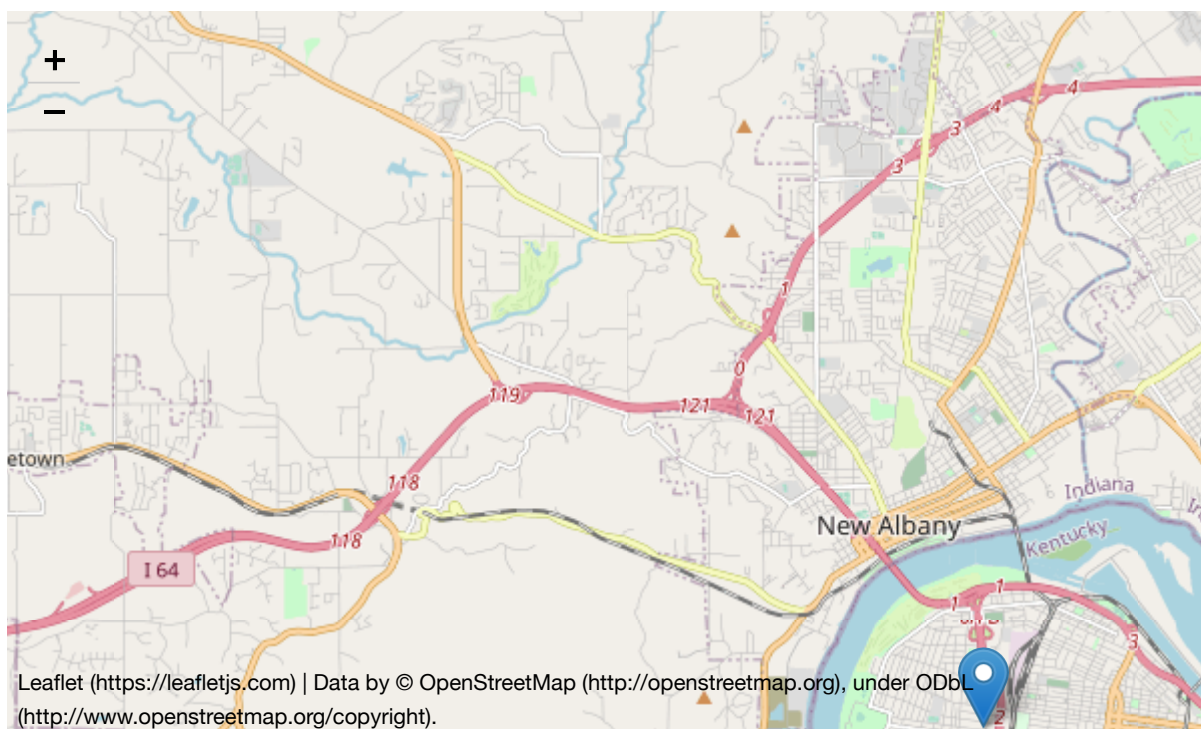
for lat, lng in zip(sample2['StartLatitude'],
                    sample2['StartLongitude']):
    folium.Marker([lat, lng], tooltip=[lat, lng]).add_to(cluster)

map_cluster

```



Out [32]:



## Heat Map of Trips

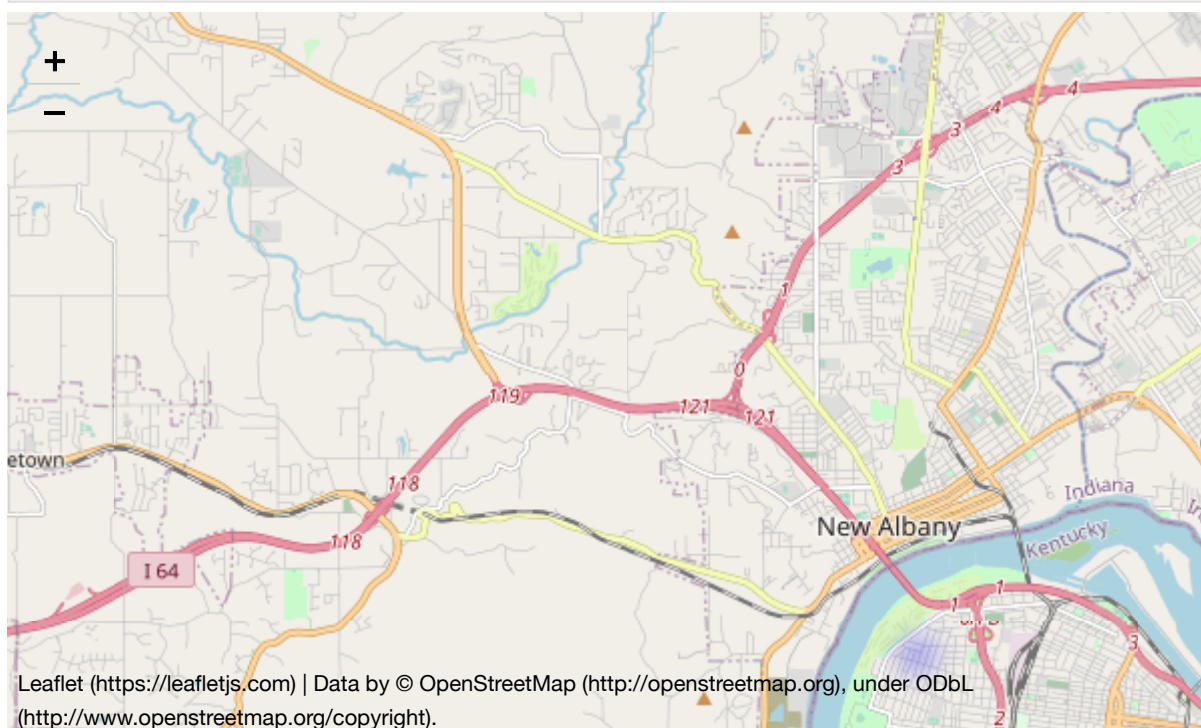
In [33]:

```
# visualise data points on a heatmap
heatmap = folium.Map(location=Louisville,
                      zoom_start=12)

HeatMap(data=sample[['StartLatitude', 'StartLongitude']].groupby(['StartLatitude',
                                                                    'StartLongitude']).agg('count').reset_index(),
        radius=8, max_zoom=12).add_to(heatmap)

heatmap
```

Out [33]:



## Map the Trips

```
In [34]: ## Map trip data (Start -> End points)
map_trip = folium.Map(location=Louisville,
                      zoom_start=12)

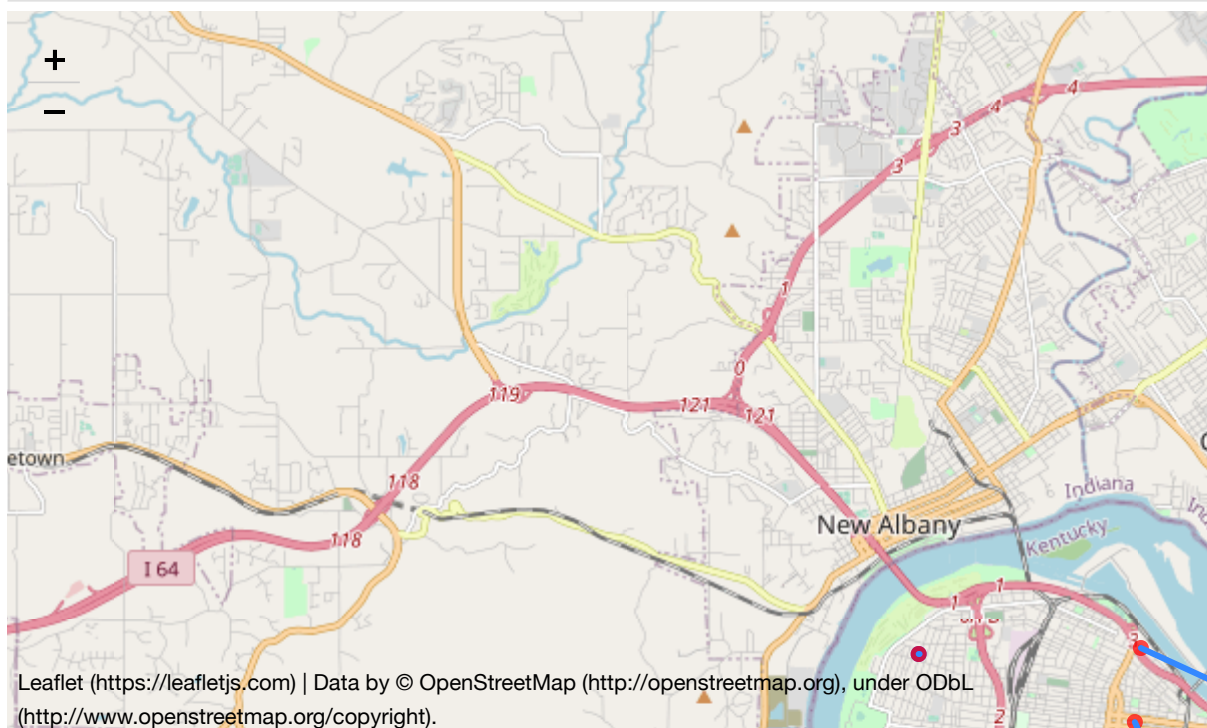
for i, row in sample.iterrows():
    folium.CircleMarker([row['StartLatitude'], row['StartLongitude']],
                        radius=4,
                        stroke=False,
                        fill_color='blue',
                        fill_opacity=0.7).add_to(map_trip)

    folium.CircleMarker([row['EndLatitude'], row['EndLongitude']],
                        radius=4,
                        stroke=False,
                        fill_color='red',
                        fill_opacity=0.7).add_to(map_trip)

    folium.PolyLine([[row['StartLatitude'], row['StartLongitude']],
                    [row['EndLatitude'], row['EndLongitude']]],
                    strokeColor= "#000000"
                    ).add_to(map_trip)

map_trip
```

Out [34]:



## Machine learning

First we will build LinerRegression Model and evaluate the MAE.

Next, we will do Feature selection using DecisionTreeRegressor and then use it to build LinerRegression model and evaluate.



Lastly, we remove outliers and build & evaluate LinerRegression model.

In [35]: `df.dtypes`

```
Out[35]: TripID                object
StartDate      datetime64[ns]
StartTime      object
EndDate        datetime64[ns]
EndTime        object
TripDuration    int64
TripDistance    float64
StartLatitude   float64
StartLongitude  float64
EndLatitude     float64
EndLongitude    float64
DayOfWeek       int64
HourNum         int64
DayOfWeekStr    object
dtype: object
```

## Linear Regression Model

```
In [36]: ## Lets predict TripDuration to TripDistance
# define X-y axis, excluding non-numerical values
y = df['TripDuration'] # dependent variable
X = df.select_dtypes(exclude=['object', 'datetime64[ns]']).drop(['TripDuration'])
```

```
In [37]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.33, random
```

```
In [38]: from sklearn import linear_model

lr = linear_model.LinearRegression()
lr_model = lr.fit(X_train, y_train)
lr_pred = lr_model.predict(X_test)
```

```
In [39]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import statistics

MAE = mean_absolute_error(y_test, lr_pred)
MSE = mean_squared_error(y_test, lr_pred)
RMSE = np.sqrt(MSE)
R2 = r2_score(y_test, lr_pred)


print("MAE : %6.2f" % (MAE))
print("MSE : %6.2f" % (MSE))
print("RMSE: %6.2f" % (RMSE))
print("R2  : %6.2f" % (R2))

MAE :    9.39
MSE :  149.80
RMSE:   12.24
R2  :    0.19
```


```
In [40]: # Visualizing model performance
plt.scatter(y_test, lr_pred)
```

```
plt.xlabel('Actual Values')
plt.ylabel('Predictions')

# Ideal predictions plot
plt.plot(y_test,y_test,'r');
```



```
In [41]: # Plotting residuals
fig = plt.figure(figsize=(10,5));
residuals = (y_test - lr_pred)
sns.histplot(residuals);
```



## DecisionTreeRegressor - Feature selection

```
In [42]: # import library
from sklearn.tree import DecisionTreeRegressor

# define the model with DecisionTreeRegressor
model = DecisionTreeRegressor()
# fit the model
model.fit(X_train, y_train)
```

Out[42]: `DecisionTreeRegressor`  
`DecisionTreeRegressor()`


```
In [43]: importance = model.feature_importances_

print(importance[0])

# summarize feature importance
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))

# plot feature importance
plt.bar([x for x in range(len(importance))], importance)
plt.show()
```

```
0.3035207405855179
Feature: 0, Score: 0.30352
Feature: 1, Score: 0.08651
Feature: 2, Score: 0.10848
Feature: 3, Score: 0.19526
Feature: 4, Score: 0.12002
Feature: 5, Score: 0.08830
Feature: 6, Score: 0.09790
```



```
In [44]: # recursive feature elimination
from sklearn.feature_selection import RFE

# define method
rfe = RFE(estimator=DecisionTreeRegressor(), n_features_to_select=2)
# fit the model
rfe.fit(X,y)
# transform the data
X_rfe = rfe.transform(X)
print("num features: %d" % rfe.n_features_)
print("selected features: %s" % rfe.support_)
print("feature ranking: %s" % rfe.ranking_)

num features: 2
selected features: [ True False False  True False False False]
feature ranking: [1 6 3 1 2 5 4]
```

```
In [45]: # train/test split the new dataset
X_train, X_test, y_train, y_test = train_test_split(X_rfe, y, test_size = 0.33,
```

```
In [46]: # linear regression on refined dataset
lr = linear_model.LinearRegression()
lr_model = lr.fit(X_train, y_train)

lr_pred2 = lr_model.predict(X_test)
```

```
In [47]: # performance metrics
MAE2 = mean_absolute_error(y_test, lr_pred2)
MSE2 = mean_squared_error(y_test, lr_pred2)
RMSE2 = np.sqrt(MSE2)
R22 = r2_score(y_test, lr_pred2)
print("MAE : %6.2f" % (MAE2))
print("MSE : %6.2f" % (MSE2))
print("RMSE: %6.2f" % (RMSE2))
print("R2 : %6.2f" % (R22))
```

```
MAE : 10.33
MSE : 169.39
RMSE: 12.24
R2 : 0.08
```

### Reducing features didn't made the predictions better.

It seems that the other variables such as origin/destination positions slightly plays a role in the trip duration too.

- another way to improve our machine learning algorithm will be to delete outliers from the dataset

## Remove Outliers

```
In [48]: ## Focus only where TripDuration is <=20min

df_20 = df[(df.TripDuration < 20)]

y_sub = df_20['TripDuration'] # dependent variable
X_sub = df_20.select_dtypes(exclude=['object', 'datetime64[ns]']).drop(['TripDur
```

```
In [49]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_sub, y_sub, test_size = 0.2)
```

```
In [50]: # linear regression
from sklearn import linear_model

lr = linear_model.LinearRegression()
lr_model = lr.fit(X_train, y_train)

lr_pred = lr_model.predict(X_test)
```


```
In [51]: # calculate statistical measures
MAE = mean_absolute_error(y_test, lr_pred)
MSE = mean_squared_error(y_test, lr_pred)
RMSE = np.sqrt(MSE)
R2 = r2_score(y_test, lr_pred)
print("MAE : %6.2f" % (MAE))
```

```
print("MSE : %6.2f" % (MSE))
print("RMSE : %6.2f" % (RMSE))
print("R2 : %6.2f" % (R2))
```


```
MAE : 3.47
MSE : 18.96
RMSE : 4.35
R2 : 0.26
```

```
In [52]: # Visualizing model performance
plt.scatter(y_test, lr_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predictions')

# Ideal predictions plot
plt.plot(y_test, y_test, 'r');
```



```
In [53]: # Plotting residuals
fig = plt.figure(figsize=(10,5))
residuals = (y_test - lr_pred)
sns.histplot(residuals);
```



# Conclusion

- We performed Model Selection and improved the model performance by doing feature selection and limiting the values and experienced the MAE reduction from **9.39 -> 3.47**

## Additional References

1. Open Mobility Foundation: [www.openmobilityfoundation.org](http://www.openmobilityfoundation.org)
2. About MDS: <https://www.openmobilityfoundation.org/about-mds/>
3. Possibilities: <https://www.openmobilityfoundation.org/whats-possible-with-mds/>
4. GIT : <https://github.com/openmobilityfoundation/mobility-data-specification>
5. SF Mobility: <https://www.sfmta.com/shared-mobility-311-complaints-location>
6. LA Dot: <https://ladot.io/wp-content/uploads/2018/12/What-is-MDS-Cities.pdf>